

Junyoung Park

june295921@cau.ac.kr | 010-2389-4303 | Seoul, South Korea

RESEARCH INTERESTS

My research interests lie in LLM Safety and Evaluation, particularly in understanding instruction-following failures and refusal dynamics during generation. I aim to develop process-level diagnostic methods that reveal when, why, and how LLMs fail or remain safe, and to connect these signals to reliable evaluation, post-training, and agent safety.

EDUCATION

Chung-Ang University (Seoul, South Korea)

Mar 2022 - Feb 2029 (Expected)

Bachelor of Art and Technology

Bachelor of Science in Cyber Security (Convergence Major)

RESEARCH EXPERIENCE

Chung-Ang University

Cyber Physical System Security Lab (PI: Jaewoo Lee)

Sep 2024 - Present


- Conducted AI/security research on LLM safety failures, refusal dynamics, multi-turn memory attacks, GraphRAG-based regulation QA, federated learning, and industrial-control anomaly detection.
- Built benchmark and evaluation pipelines for collecting generations, extracting logit-level safety signals, running LLM-as-Judge checks, and producing reproducible metric tables and figures.
- Connected research prototypes to external outputs, including conference presentation material, software/copyright registration, and industry-academic research documentation.

SELECTED RESEARCH OUTPUTS

Beyond Attack Success Rate: Temporal Logit Observability for LLM Safety Failures

First Author | *arXiv:2605.29629*

- Diagnoses safety failure formation by tracking refusal/compliance margins and harmful-token tendencies across decoding steps, beyond outcome-only ASR.



(a) Outcome Metrics Hide the Decoding Process


(b) Logit Signals Make Temporal Dynamics Observable

Outcome-only ASR can hide different decoding-process dynamics; TLO makes generation-time safety signals observable.

Persona Attack: Incremental Memory Injection Jailbreak Attack against Large Language Models

First Author | *Manuscript*

- Analyzes instruction-following failures and refusal dynamics under adversarial multi-turn memory contexts.



Incremental memory injection accumulates attacker-controlled context across turns before inserting a new harmful request.

A GraphRAG Framework for Financial Security Regulation

Co-author | 2025 Fall Conference of Society for e-Business Studies, Oral

- Built a relation-aware regulation QA pipeline using E5-ko semantic search, FAISS indexing, NetworkX graph expansion, and evidence blocks.

Software Registration & Copyright

Sep 2025 - Nov 2025

- GraphRAG: Korea Copyright Commission (C-2025-062232) and NIPA (ASSET_0014943); Distributed Unlearnable Example: Korea Copyright Commission (C-2025-051859).

PROJECTS

Fairness-Aware Machine Unlearning for Face Landmark Detection

Mar 2025 - Jun 2025

- Built a PyTorch CNN eye-localization model and analyzed group-wise fairness degradation after unlearning Asian face samples.




Visual comparison of original, unlearned, and refined landmark predictions in the SafeAI project.

Industrial Control System Anomaly Detection on the HAI Dataset

Sep 2024 - Dec 2024

- Developed a BiLSTM autoencoder-based time-series anomaly detection pipeline with windowed modeling and TaPR-oriented threshold search.



HAI process structure used for industrial-control time-series anomaly detection experiments.

On-device Machine Learning App

Jan 2024 - Jun 2024

- Developed a Swift-based iOS diary app that performs on-device emotion analysis with Create ML, stores records via HealthKit, and visualizes mood trends.

TECHNICAL SKILLS

Python, PyTorch, Transformers, pandas, NumPy, FAISS, NetworkX, LLM-as-Judge evaluation, benchmark automation, Swift, Create ML, HealthKit, Git.